

Threshold true prevalence:

Model	TP	FN	FP	TN	Sensitivity/Recall	Specificity	PPV/Precision	NPV	F1
LR-Kelly	3,020	17,562	12,312	1,735,885	14.67%	99.30%	19.70%	99.00%	16.82%
LR-Rick	3,913	16,669	17,830	1,730,367	19.01%	98.98%	18.00%	99.05%	18.49%
LR-Full List 299	3,464	17,118	14,851	1,733,346	16.83%	99.15%	18.91%	99.02%	17.81%
Transformer-Kelly	4,525	16,057	17,112	1,731,085	21.99%	99.02%	20.91%	99.08%	21.44%
Transformer-Rick	4,726	15,856	15,151	1,733,046	22.96%	99.13%	23.78%	99.09%	23.36%
Transformer-Full List 299	4,654	15,928	14,305	1,733,892	22.61%	99.18%	24.55%	99.09%	23.54%

Values are exact confusion-matrix counts on the held-out test split (exp_calib_july), evaluated at the real cohort prevalence. Threshold is the isotonic-calibrated prevalence_matched cutoff, selected on the calibration split only (never the test split). The test split contains 1,768,779 patients (single train/calibration/test split, not 10-fold cross-validation, so counts are not fold averages).